

# Investigating discrimination bias in predictive modelling

Sara Thiringer   Jonathan Rittmo

26/10/2020

# Background: The Problem

In recent years, many scandalous examples have shown that statistical models trained on large amounts of data can “act” discriminatory. Examples include:

- ▶ Adds of high-income jobs being shown less frequently to women, presumable because they've been predicted to be less interested or suitable<sup>1</sup>
- ▶ Black people's health status being underestimated, leading to inappropriate health care measures<sup>2</sup>
- ▶ Black people being predicted a higher risk for crime recidivism, leading to higher penalties<sup>3</sup>

---

<sup>1</sup>Datta, Tschantz, and Datta (2015)

<sup>2</sup>Obermeyer et al. (2019)

<sup>3</sup>ProPublica (2016)

# Project Aims

- ▶ How can we quantify fairness in order to be able to evaluate algorithmic fairness?
- ▶ What methods are available to increase algorithmic fairness?  
In what type of situations do they apply? (i.e. In what kind of situations can we expect them to be successful?)

# Background: Why Discrimination Bias?

- ▶ Correlation between outcome  $y$  and protected characteristic  $x_p$
- ▶ Correlation between important predictors  $x_i$  and protected characteristic  $x_p$
- ▶ Under/over sampling of groups with protected characteristic  $x_p$

## Possible Solutions

<b>Pre-Processing</b>	<b>Training</b>	<b>Prediction</b>
Resampling	Penalty	Threshold
Mapping	Model bias	adjustments
Altering labels	Tuning for fairness	Alter predictions

We've chosen to work with resampling and threshold adjustment.

## Possible Goals

Demographic parity

$$P(Y = 1|X = 1) = P(Y = 1|X = 0)$$

Equalized odds

$$P(G = 1|X = 0, Y = 1) = P(G = 1|X = 1, Y = 1)$$

Data Frame Summary  
 COMPAS  
 Dimensions: 5278 x 7  
 Duplicates: 4706





No	Variable	Stats / Values	Freqs (% of Valid)
1	Two_yr_Recidivism [factor]	1. 0 2. 1	2795 (53.0%) 2483 (47.0%)
2	Number_of_Priors [integer]	Mean (sd) : 3.5 (4.9) min < med < max: 0 < 2 < 38 IQR (CV) : 5 (1.4)	36 distinct values
3	Above45 [factor]	1. 0 2. 1	4182 (79.2%) 1096 (20.8%)
4	Below25 [factor]	1. 0 2. 1	4122 (78.1%) 1156 (21.9%)
5	Misdemeanor [factor]	1. 0 2. 1	3440 (65.2%) 1838 (34.8%)
6	Ethnicity [factor]	1. African_American 2. Caucasian	3175 (60.2%) 2103 (39.8%)
7	Sex [factor]	1. Female 2. Male	1031 (19.5%) 4247 (80.5%)

# Data

Variable	Type	Values
Two_yr_Recidivism	Factor	1 / 0
Number_of_Priors	Numerical	Mean (sd) : 3.5 (4.9)
Above45	Factor	1 / 0
Below25	Factor	1 / 0
Misdemeanor	Factor	1 / 0
Ethnicity	Factor	African_American / Caucasian
Sex	Factor	Female / Male



# Descriptives




# Models

Model	Tuning
Random Forest	Predictors at each split
Artificial neural net	Number of hidden nodes
Logistic ridge regression	Penalisation
K-nearest neighbour (left out)	Number of neighbours
AdaBoost	Predictors at each split

# Accuracy and Fairness for the Initial Models

```
## Warning: namespace 'xgboost' is not available and has been
## by .GlobalEnv when processing object 'fobject1'
```




# Disparate Impact Removal

Removes differences between groups while preserving it within the groups. Example: Number of priors.

```
## Warning: Removed 71 rows containing non-finite values (s
```

```
## Warning: Removed 71 rows containing non-finite values (s
```



## Resampling

Using undersampling and oversampling to even out inequalities between the deprived and privileged groups having positive and negative outcome attributes respectively.

Table 5: Joint distribution of Ethnicity and Recidivism

	African_American	Caucasian
0	1514	1281
1	1661	822

Note: 1 = Recidivism

# Uniform Resampling

Aim: Make the joint distribution of Ethnicity and Two\_yr\_Recidivism uniform by duplicating some observations and removing others.

Table 6: Joint distribution of Ethnicity and Recidivism, uniform resampling

	African_American	Caucasian
0	1352	885
1	1201	786

Note: 1 = Recidivism

# Preferential Resampling

- ▶ Very similar to uniform
- ▶ Borderline observations skipped or duplicated more often
- ▶ Evaluated by logistic regression


Table 7: Joint distribution of Ethnicity and Recidivism, preferential resampling

	African_American	Caucasian
0	1352	885
1	1201	786

Note: 1 = Recidivism

# Comparison

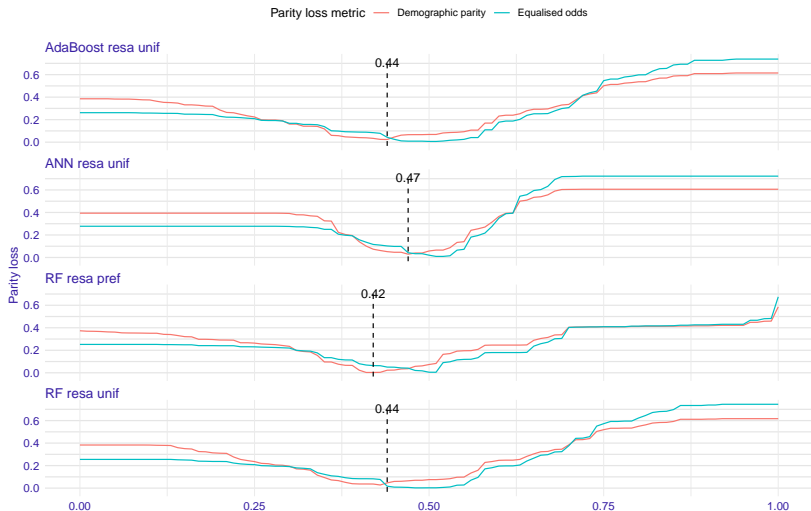
```
## Warning: namespace 'xgboost' is not available and has been
## by .GlobalEnv when processing object 'fobject_all'
```





# Threshold Adjustment

```
## Warning: namespace 'xgboost' is not available and has been
## by .GlobalEnv when processing object 'fobject_co'
```



# Comparing Models with Different Thresholds

```
## Warning: namespace 'xgboost' is not available and has been
## by .GlobalEnv when processing object 'fo_comp'
```



# Evaluation on Test Data - Comparison

```
## Warning: namespace 'xgboost' is not available and has been
## by .GlobalEnv when processing object 'fobject_test'
```



## Final Model Performance and Conclusion

- ▶ The resampling methods seems to have best effect on minimising parity loss while also preserving accuracy rates
- ▶ Random Forest looked most promising, but while tested on new data Logistic Regression, Artifical Neural Network and AdaBoost were equally strong competitors

Table 8: Performance

recall	precision	accuracy	auc
0.717	0.654	0.649	0.7

# General Conclusions

- ▶ It is possible to build models who satisfy some fairness criteria without a too large drop in accuracy
- ▶ Decisions for fair models include: fairness measure, evaluation metrics and choice of methods
- ▶ Less complex methods can be found amongst resampling and threshold adjustment
- ▶ As always, what type of data we are dealing with will largely impact the results of different methods (for example why disparate impact remover didn't work)

## References

- Datta, Amit, Michael Carl Tschantz, and Anupam Datta. 2015. "Automated Experiments on Ad Privacy Settings." *Proceedings on Privacy Enhancing Technologies* 2015 (1): 92–112.  
<https://doi.org/https://doi.org/10.1515/popets-2015-0007>.
- Obermeyer, Ziad, Brian Powers, Christine Vogeli, and Sendhil Mullainathan. 2019. "Dissecting Racial Bias in an Algorithm Used to Manage the Health of Populations." *Science* 366 (6464): 447.  
<https://doi.org/10.1126/science.aax2342>.